

End-to-end success by condition: the omni pathway against the reference cascade

DESCRIPTIVE (H4a). 31,616 rows: every omni row plus only the Whisper-Turbo -> large-router cascade rows, so the omni model is not being compared against a weaker cascade. Whiskers are a 95% cluster bootstrap over the 416 utterances (B = 2000, seed 42). NOTE: The two pathways differ in capacity as well as in architecture; the confirmatory statement is the interaction in fig4, which cancels a constant capacity gap. NOTE: Grey crosses mark conditions where dropping the 748 `full\_hallucination` rows moves the cascade mean — the omni advantage under babble is carried by invented Whisper text (§13).

